

# **Debugging a C program**

Lab 5

# Debugging

- Finding and eliminating errors from programs
- Grace Hopper and the “First actual case of bug being found”



9/9

0800 Antan started  
 1000 " stopped - antan ✓

1300 (032) MP-MC { 1.2700 9.037 847 025  
 2.130476415 } 9.037 846 895 convd  
 (033) PRO 2 2.130476415  
 convd 2.130676415

Relays 6-2 in 033 failed speed test  
 in relay 11.00 test.

Relays changed  
 1100 Started Cosine Tape (Sine check)  
 1525 Started Multi Adder Test.

1545 Relay #70 Panel F  
 (moth) in relay.

First actual case of bug being found.  
 1630 Antan started.  
 1700 closed down.

Relay 3145  
 Relay 3376

# Debugging Process

- Reproduce the bug
- Simplify program input
- Use a debugger to track down the origin of the problem
- Fix the problem

# Debugger

- A program that is used to run and debug other (target) programs
- Advantages:
  - Programmer can:
    - step through source code line by line
      - each line is executed on demand
    - interact with and inspect program at run-time
    - If program crashes, the debugger outputs where and why it crashed

# GDB – GNU Debugger

- Debugger for several languages
  - C, C++, Java, Objective-C
- Allows you to inspect what the program is doing at a certain point during execution
- Logical errors and segmentation faults are easier to find with the help of gdb

# Using GDB

## 1. Compile Program

- Normally: `gcc [flags] <source files> -o <output file>`
- Debugging: `gcc [other flags] -g <source files> -o <output file>`
  - enables built-in debugging support

## 2. Specify Program to Debug

- `$ gdb <executable>`
- or
- `$ gdb`
- `(gdb) file <executable>`

# Using GDB

## 4. Run Program

- `(gdb) run`                      or
- `(gdb) run [arguments]`

## 5. In GDB Interactive Shell

- Tab to Autocomplete, up-down arrows to recall history
- `help [command]` to get more info about a command

## 6. Exit the gdb Debugger

- `(gdb) quit`

# Run-Time Errors

- Segmentation fault
  - Program received signal SIGSEGV, Segmentation fault.  
0x0000000000400524 in *function* (arr=0x7fffc902a270, r1=2, c1=5, r2=4, c2=6) at *file.c*:12
    - Line number where it crashed and parameters to the function that caused the error
- Logic Error
  - Program will run and exit successfully
- How do we find bugs?



# Setting Breakpoints

- Breakpoints
  - used to stop the running program at a specific point
  - If the program reaches that location when running, it will pause and prompt you for another command
- Example:
  - (gdb) `break file1.c:6`
    - Program will pause when it reaches line 6 of file1.c
  - (gdb) `break my_function`
    - Program will pause at the first line of `my_function` every time it is called
  - (gdb) `break [position] if expression`
    - Program will pause at specified position only when the expression evaluates to true

# Breakpoints

- Setting a breakpoint and running the program will stop program where you tell it to
- You can set as many breakpoints as you want
  - (gdb) `info breakpoints/break/br/b`  
shows a list of all breakpoints

# Deleting, Disabling and Ignoring BPs

- (gdb) delete [bp\_number | range]
  - Deletes the specified breakpoint or range of breakpoints
- (gdb) disable [ *bp\_number* | *range* ]
  - Temporarily deactivates a breakpoint or a range of breakpoints
- (gdb) enable [ *bp\_number* | *range* ]
  - Restores disabled breakpoints
- If no arguments are provided to the above commands, all breakpoints are affected!!
- (gdb) ignore *bp\_number iterations*
  - Instructs GDB to pass over a breakpoint without stopping a certain number of times.
    - bp\_number: the number of a breakpoint
    - Iterations: the number of times you want it to be passed over

# Displaying Data

- Why would we want to interrupt execution?
  - to see data of interest at run-time:
  - (gdb) `print [/format] expression`
    - Prints the value of the specified expression in the specified format
  - Formats:
    - d: Decimal notation (default format for integers)
    - x: Hexadecimal notation
    - o: Octal notation
    - t: Binary notation

# Resuming Execution After a Break

- When a program stops at a breakpoint
  - 4 possible kinds of gdb operations:
    - **c or continue**: debugger will continue executing until next breakpoint
    - **n or next**: debugger will execute the next line as single instruction
    - **s or step**: same as next but functions are executed line by line instead of as a single instruction
    - **f or finish**: debugger will resume execution until the current function returns. Execution stops immediately after the program flow returns to the function's caller
      - the function's return value and the line containing the next statement are displayed

# Watchpoints

- Watch/observe changes to variables
  - (gdb) `watch my_var`
    - sets a watchpoint on `my_var`
    - the debugger will stop the program when the value of *my\_var* changes
    - old and new values will be printed
  - (gdb) `rwatch expression`
    - The debugger stops the program whenever the program reads the value of any object involved in the evaluation of *expression*

# Process Memory Layout

Args and env vars

Command line arguments  
and environment variables

Stack  
|  
V

Grows downward as a  
result of a newly called  
function

Unused memory

^  
|  
Heap

grows upward when dynamic  
memory is requested by C's  
malloc() or C++'s new

Uninitialized Data Segment (bss)

Uninitialized global variables

Initialized Data Segment

Initialized global variables

Text Segment

code to be executed

# Stack Info

- A program is made up of one or more functions which interact by calling each other
- For each function call, an area of memory is set aside. This area of memory is called a **stack frame**
- A stack frame holds the following crucial info:
  - storage space for all the local variables
  - the memory address to return to when the called function returns
  - the arguments to the called function
- Each function call gets its own stack frame. Collectively, all the stack frames make up the **call stack**



# Stack Frames and the Stack

```

1  #include <stdio.h>
2  void first_function(void);
3  void second_function(int);
4
5  int main(void)
6  {
7      printf("hello world\n");
8      first_function();
9      printf("goodbye goodbye\n");
10
11     return 0;
12 }
13
14
15 void first_function(void)
16 {
17     int imidate = 3;
18     char broiled = 'c';
19     void *where_prohibited = NULL;
20
21     second_function(imidate);
22     imidate = 10;
23 }
24
25
26 void second_function(int a)
27 {
28     int b = a;
29 }

```

Frame for `main()`

Frame for `first_function()`

Return to `main()`, line 9

Storage space for an int

Storage space for a char

Storage space for a void \*

Frame for `second_function()`:

Return to `first_function()`, line 22

Storage space for an int

Storage for the int parameter named `a`

When `first_function()` is called from `main()`, it is used to determine where to put the frame for `main()`. It holds information about the frame for `main()`, including the address of the frame for `main()`, the address of the frame for `first_function()`, and the address of the frame for `second_function()`. The frame for `main()` is used to store the return address of `main()`, the address of the frame for `first_function()`, and the address of the frame for `second_function()`. The frame for `first_function()` is used to store the return address of `first_function()`, the address of the frame for `main()`, and the address of the frame for `second_function()`. The frame for `second_function()` is used to store the return address of `second_function()`, the address of the frame for `main()`, and the address of the frame for `first_function()`.

# Analyzing the Stack in GDB

- `(gdb) backtrace/bt`
  - Shows the call trace (the call stack)
  - Without function calls:
    - #0 main () at program.c:10
    - one frame on the stack, numbered 0, and it belongs to main()
  - After call to function `display()`
    - #0 display (z=5, zptr=0xbffffb34) at program.c:15
    - #1 0x08048455 in main () at program.c:10
    - Two stack frames: frame 1 belonging to main() and frame 0 belonging to display().
    - Each frame listing gives
      - the arguments to that function
      - the line number that's currently being executed within that frame

# Analyzing the Stack

- (gdb) info frame
  - Displays information about the current stack frame, including its return address and saved register values
- (gdb) info locals
  - Lists the local variables of the function corresponding to the stack frame, with their current values
- (gdb) info args
  - List the argument values of the corresponding function call

# Other Useful Commands

- (gdb) info functions
  - Lists all functions in the program
- (gdb) list
  - Lists source code lines around the current line

# Lab 5

- Download old version of coreutils with buggy ls program
  - Untar, configure, make
- Bug: ls -lt mishandles files whose time stamps are very far in the past. It seems to act as if they are in the future

```
$ touch -d '1918-11-11 11:00 GMT' wwi-armistice
$ touch now
$ sleep 1
$ touch now1
$ ls -lt wwi-armistice now now1
```

Output:

```
-rw-r--r-- 1 eggert eggert 0 Nov 11 1918 wwi-armistice
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now1
-rw-r--r-- 1 eggert eggert 0 Feb 5 15:57 now
```

# Goal: Fix the Bug

- **Reproduce the Bug**
  - Follow steps on lab web page
- **Simplify input**
  - Run ls with -l and -t options only
- **Debug**
  - Use gdb to figure out what's wrong
  - \$ gdb ./ls
  - (gdb) run -lt wwi-armistice now now1  
(run from the directory where the compiled ls lives)
- **Patch**
  - Construct a patch "lab5.diff" containing your fix
  - It should contain a ChangeLog entry followed by the output of diff -u

# Lab Hints

- Don't forget to answer all questions! (lab5.txt)
- Make sure you don't submit a reverse patch! (lab5.diff)
- "Try to reproduce the problem in your home directory, instead of the \$tmp directory. How well does SEASnet do?"
  - Timestamps represented as seconds since Unix Epoch
  - SEASnet NFS filesystem has unsigned 32-bit time stamps
  - Local File System on Linux server has signed 32-bit time stamps
  - If you touch the files on the NFS filesystem it will return timestamp around 2054
  - => files have to be touched on local filesystem (df -l)
- Use "info functions" to look for relevant starting point
- Compiler optimizations: -O2 -> -O0
  - ./configure CGLAGS="...-O0"