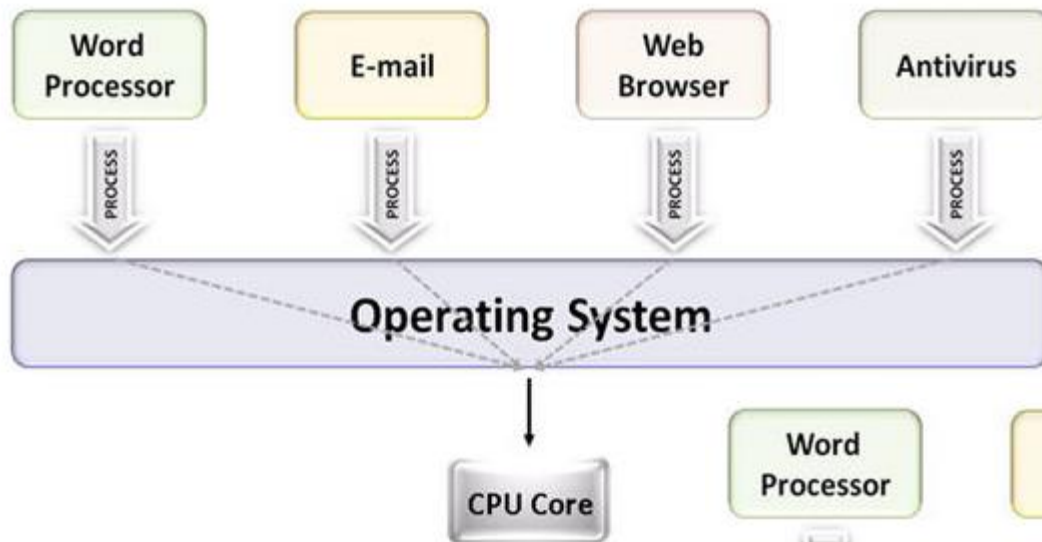


Multithreaded Performance

Week 8

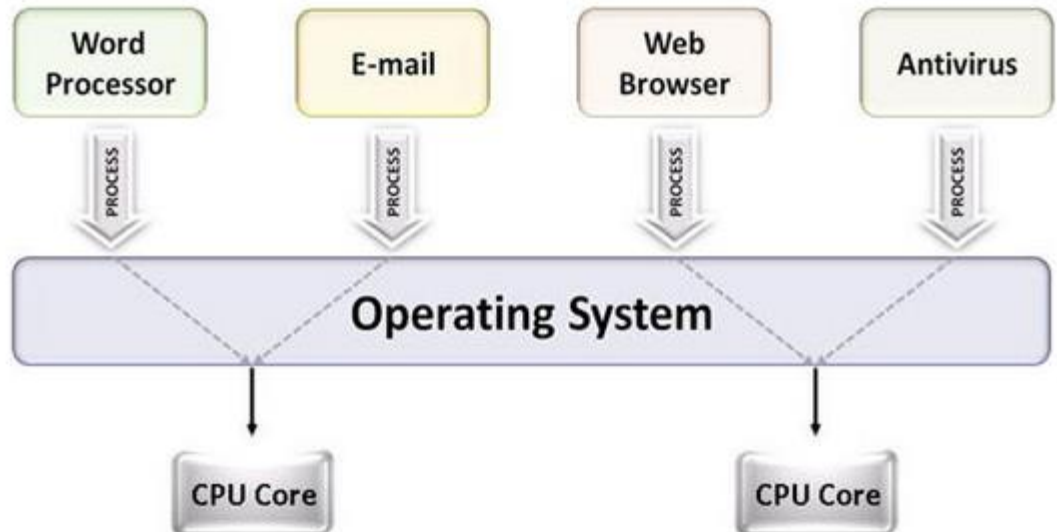
Multiprocessing

- The use of multiple CPUs/cores to run multiple tasks simultaneously



**Uniprocessing
system**

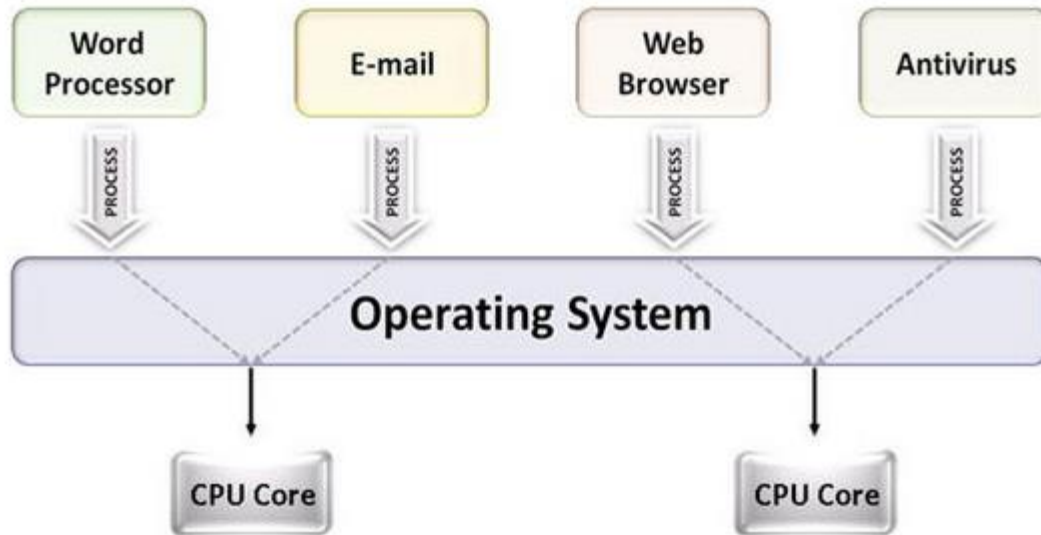
**Multiprocessing
system**



Parallelism

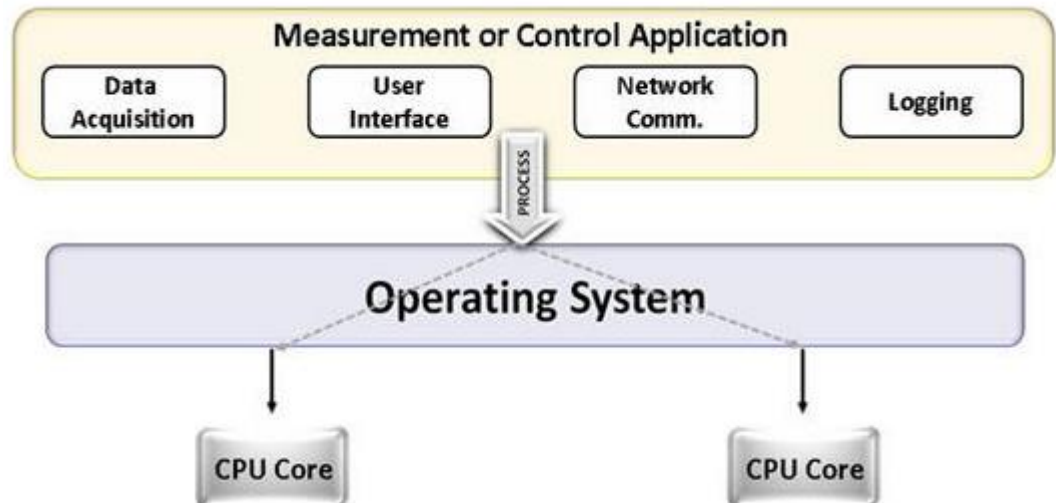
- Executing several computations simultaneously to gain performance
- Different forms of parallelism
 - **Multitasking**
 - Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
 - **Multithreading**
 - Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system

Multitasking vs. Multithreading




Multitasking

Multithreading



Multithreading


Multitasking



What is a thread?

- A flow of instructions, path of execution within a process
- The smallest unit of processing scheduled by OS
- A process consists of at least one thread
- Multiple threads can be run on:
 - **A uniprocessor (time-sharing)**
 - Processor switches between different threads
 - Virtual parallelism (concurrency)
 - **A multiprocessor**
 - Multiple processors or cores run the threads at the same time
 - True parallelism

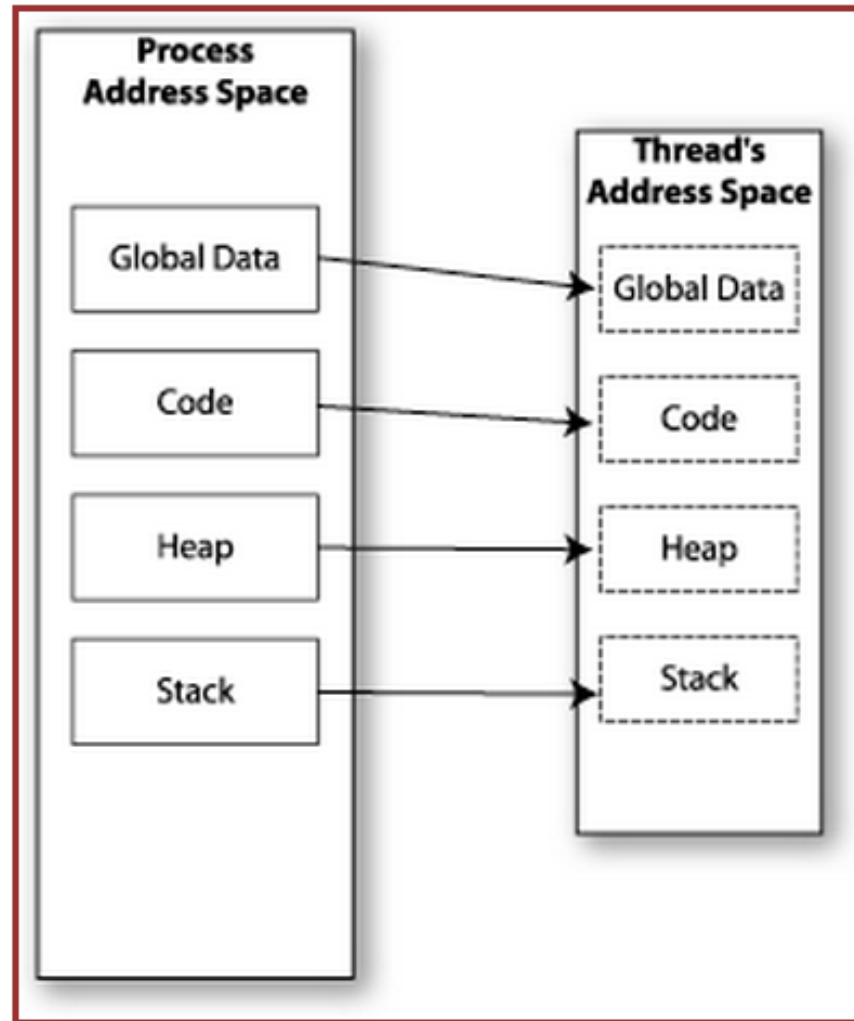
Multiple threads sharing a single CPU



Multiple threads on multiple CPUs



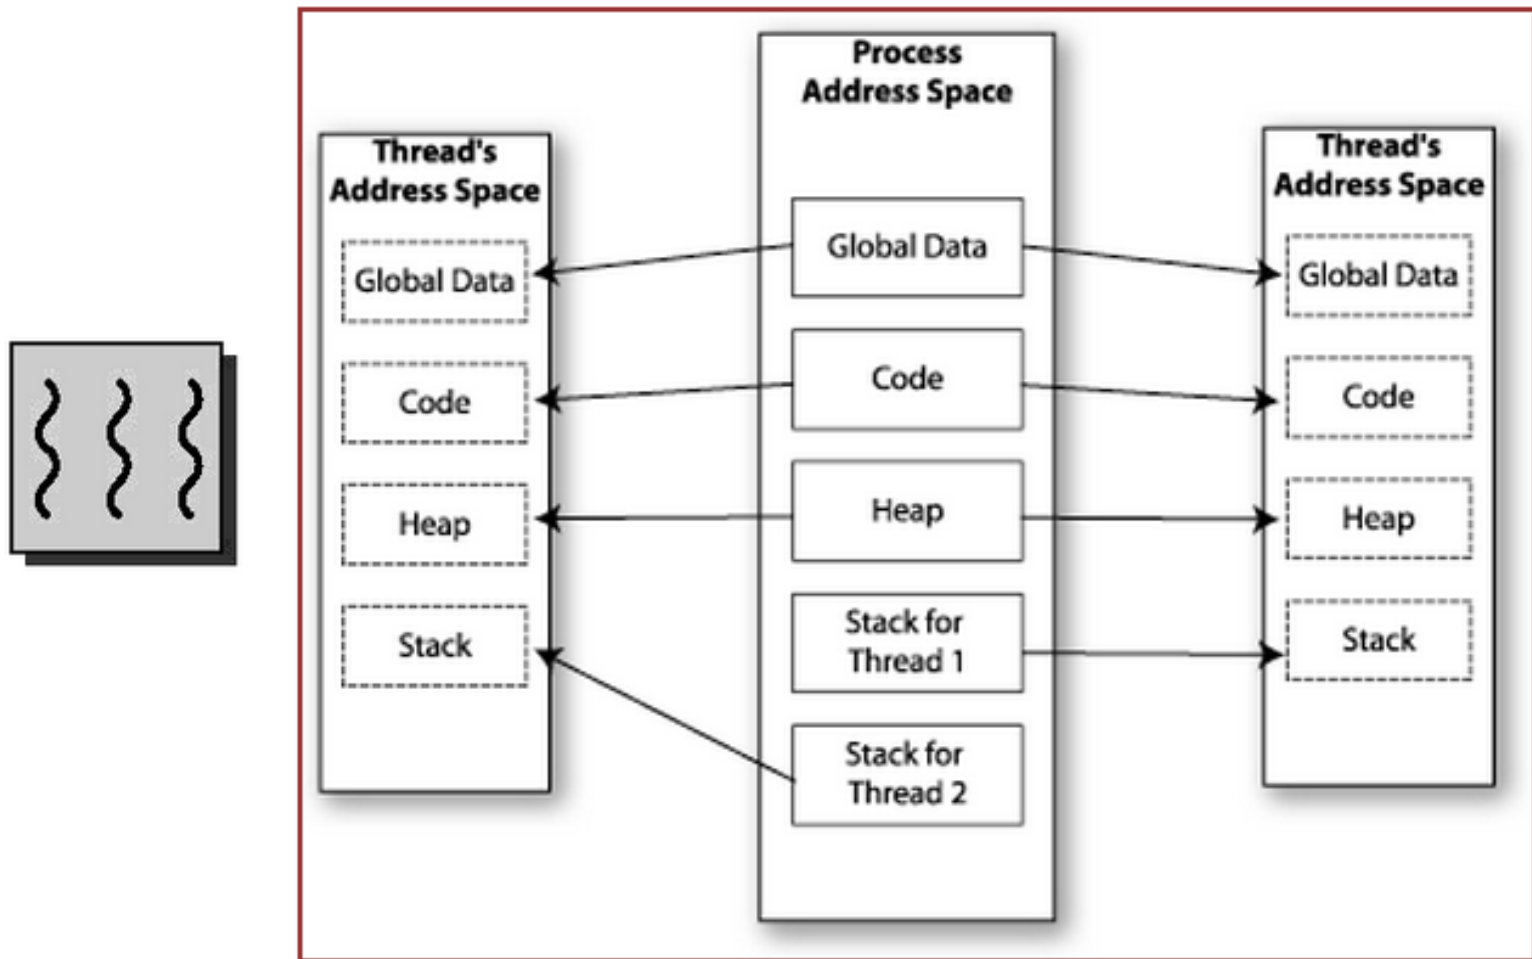
Memory Layout: Single-Threaded Program



Multitasking

- `$ tr -cs 'A-Za-z' '[\n*]' | sort -u | comm -23 – words`
 - Process 1 (tr)
 - Process 2 (sort)
 - Process 3 (comm)
- Each process has its own address space
- How do these processes communicate?
 - Pipes/System Calls

Memory Layout: Multithreaded Program



Multithreading

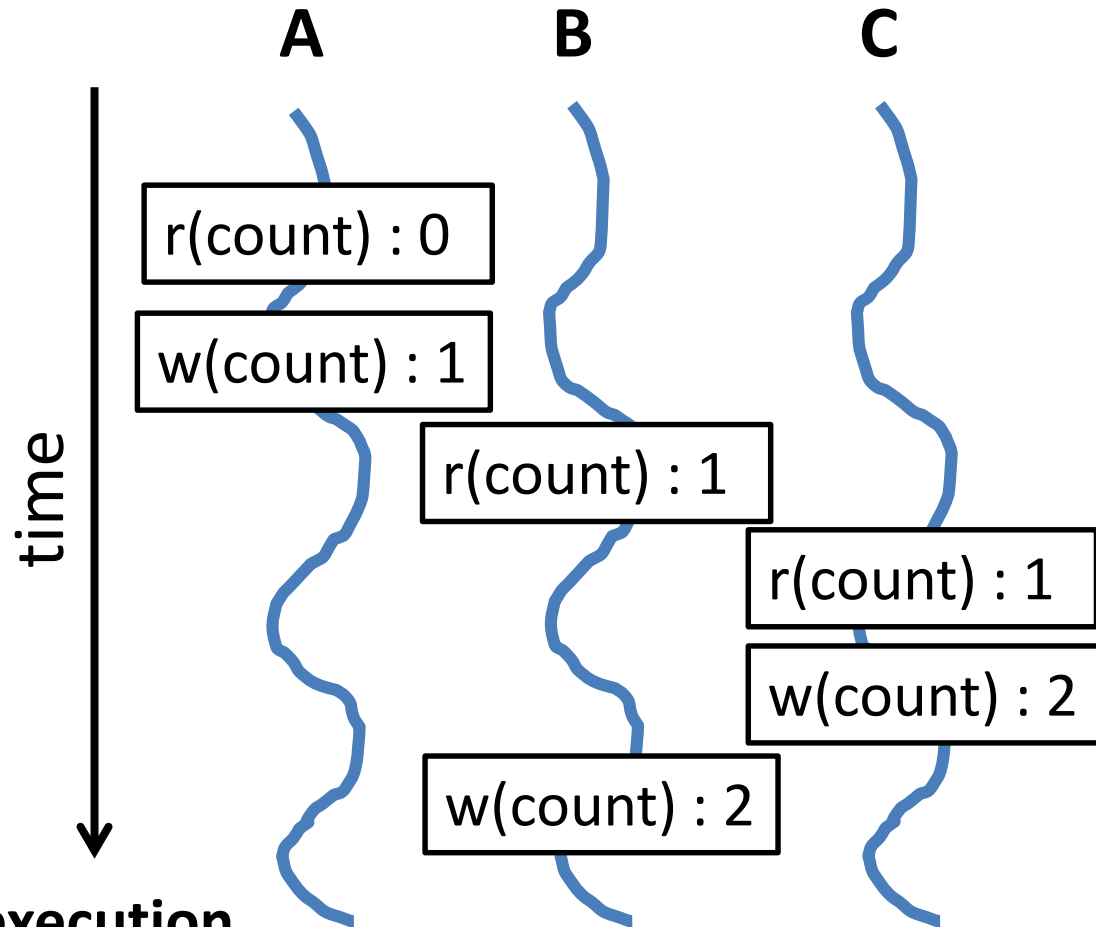
- Threads share all of the process's memory except for their stacks
- => Data sharing requires no extra work (no system calls, pipes, etc.)

Shared Memory

- Makes multithreaded programming
 - **Powerful**
 - can easily access data and share it among threads
 - **More efficient**
 - No need for system calls when sharing data
 - Thread creation and destruction less expensive than process creation and destruction
 - **Non-trivial**
 - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```



Result depends on order of execution
=> Synchronization needed

Multithreading & Multitasking: Comparison

- **Multithreading**

- Threads share the same address space
 - Light-weight creation/destruction
 - Easy inter-thread communication
 - An error in one thread can bring down all threads in process

- **Multitasking**

- Processes are insulated from each other
 - Expensive creation/destruction
 - Expensive inter-process communication (IPC)
 - An error in one process cannot bring down another process

Lab 8

- Evaluate the performance of multithreaded `sort`
- Add `/usr/local/cs/bin` to `PATH`
 - `$ export PATH=/usr/local/cs/bin:$PATH`
- Generate a file containing 10M random **double-precision floating point numbers**, one per line with no white space
 - `/dev/urandom`: pseudo-random byte generator

Lab 8

- `od`
 - write the contents of its input files to standard output in a user-specified format
 - Options
 - `-t f`: Double-precision floating point
 - `-N <count>`: Format no more than *count* bytes of input
- `sed, tr`
 - Remove address, delete spaces, add newlines between each float

Lab 8

- Use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the `-g` option: general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g --parallel=[1, 2, 4, 8] 10M_file > /dev/null`
 - Record the times and steps in `log.txt`