

Week 07M: Multithreaded Programming

Thuy Vu

- Assignment 6
 - web.cs.ucla.edu/classes/winter17/cs35L/assign/assign6.html
 - Time due 23:55 this Friday, February 24


Week 7: Assignment #6 – Multithreaded Performance

Week 8: Assignment #7 – Dynamic Linking


Week 9: Assignment #8 – SSH Setup and Use in Applications

Week 10: Assignment #9 – Change management

Parallelism is to execute several computations simultaneously




Uni-processing




Multi-processing

Parallelism is to execute several computations simultaneously



Multitasking – several processes are scheduled alternately or possibly simultaneously on a multiprocessing system




Multithreading – same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system

- **Process** – each process has its own address space


```
tr 'A-Z' 'a-z' | sort | comm -23 - f
```

- Process 1: tr
- Process 2: sort
- Process 3: comm



- **Thread** – flow of instructions in a process
 - the smallest unit of processing scheduled by OS
 - a process = a main thread + other threads
 - multi threads can be running on:

Multiple threads sharing a single CPU



one processor – switching between different threads → **pseudo** parallelism

Multiple threads on multiple CPUs



multi processors – threads run on different cores → **true** parallelism

- all threads share the same address space, but have private stacks

shared memory in threads

- **powerful** – easily access and share among threads
- **efficient** – easy to spawn/kill, no need system calls when sharing data
- **non-trivial** – race condition → needs synchronization

thread APIs include

- `pthread_create` – creates a new thread within a process
- `pthread_join` – waits for another thread to terminate
- `pthread_equal` – compares thread ids
- `pthread_self` – returns the id of the calling thread
- `pthread_exit` – terminates the currently running thread

the lab

① check version

- `sort --version`
- `export PATH=/usr/local/cs/bin/:$PATH`

② generate file of 10M random double-precision floating point numbers

- `/dev/urandom`
- **od** – to write an input file to standard output in format
 - `-t f` – double-precision floating point
 - `-N <count>` – format no more than `<count>` bytes of input
- **sed, tr** – clean up

③ `time -p` on command `sort -g` and send output to `/dev/null`

- `-g` ~ general numeric values
- `--parallel=...` ~ n-threads

the homework

- `int pthread_create(pthread_t *thread,
 const pthread_attr_t *attr,
 void *(*start_routine)(void*),
 void *arg);`
 - create a new thread with handle `thread` with attribute `attr` to execute `start_routine` with `arg` as its sole argument
- `int pthread_join(pthread_t thread, void **value_ptr);`
 - suspend execution of `the calling thread` until the target thread terminates, `unless the target thread has already terminated`
 - what can go wrong without calling `pthread_join`?


...

1 Ray Tracing

- a technique to render lights in 3D
- light rays to bounce from object to object



with-out ray-tracing



with ray-tracing

→highly computationally intensive

2 write a multi-threaded version!

