

# CS 35L

Week 7

TA: Tomer Weiss  
Feb-18-2016

goo.gl/2XwK2R

Slides

# Announcements

- Student presentations today:
  - Jai - Github
  - Kyle - What happens when facial recognition tools are available to everyone
  - Ruiyi - Moore's Law Goes Post-CMOS

[web.cs.ucla.edu/classes/winter16/cs35L/assign/assign10.html](http://web.cs.ucla.edu/classes/winter16/cs35L/assign/assign10.html)

- Next week:
  - Write your topic [here](#)
  - Not registering your topic beforehand may result in rescheduling of your presentation
  - For reference on presentation, grading, please refer to this [rubric](#).

# System Call Programming

Week 7

# System calls and Library calls usage

- System calls
  - executed by the operating system
  - perform simple single operations
- Library calls
  - executed in the user program
  - may perform several task
  - may call system calls

# System calls vs library call conventions

- Library functions often return pointers
  - `FILE *fp = fopen("cs35l","r")`
  - `NULL` for return for failure
- System calls usually return an integer
  - `int res=system_call_function(a_few_args)`
  - Where the return value
    - `res >= 0` → all is well
    - `res < 0` → failure
    - See the global variable `errno` for more info

# Reminder of how System calls work

1. program get to the system call in the user's code

```
int res = sys_call(a_few_params)
```

2. puts the parameters on the stack
3. performs a system 'trap' -- hardware switch

\*\*\*now in system mode\*\*\*

4. operating system code may copy large data structures into system memory
5. starts operation...
6. operation complete!
7. if necessary copies result data structures back to user program's memory

\*\*\*return to user mode\*\*\*

8. user program puts return code into `res`(the return value from the system call)
9. program recommences


# System calls

- `int res = ssize_t read(int fildes, void *buf, size_t nbyte)`
  - fildes: file descriptor
  - buf: buffer to write to, is not NULL terminated
  - nbyte: number of bytes to read
  - res: is 0 at end of file, negative for error
- `int res = ssize_t write(int fildes, const void *buf, size_t nbyte)`
  - fildes: file descriptor
  - buf: buffer to write to, need not be NULL terminated
  - nbyte: number of bytes to write
    - res might be less than nbyte if OS buffers full
    - should check and repeat until all gone
  - res: is 0 for end of file, negative for error



# *Input\Output*

- each running program has numbered inputs/outputs:
  - 0 standard input
    - often used as input if no file is given
    - default input from the user terminal
  - 1 standard output
    - simple program's output goes here
    - default output to user terminal
  - 2 standard error
    - error messages from user
    - default output to the user terminal



- these numbers are called **file descriptors**
  - used by system call to refer to files

## *More examples: System calls*

- `int open(const char *pathname,int flags,mode_t mode)`
- `int close(int fd)`
- File descriptors:
  - 0 stdin
  - 1 stdout
  - 2 stderr
- `pid_t getpid(void)`
  - returns the process id of the calling process
- `int dup(int fd)`
  - Duplicates a file descriptor fd. Returns a second file descriptor that points to the same file table entry as fd does.
- `int fstat(int fildes, struct stat *buf)`
  - Returns information about the file with the descriptor fildes to buf

## *More examples: System calls*

```
struct stat {
dev_t      st_dev;          /* ID of device containing file */
ino_t      st_ino;          /* inode number */
mode_t     st_mode;         /* protection */
nlink_t    st_nlink;        /* number of hard links */
uid_t      st_uid;          /* user ID of owner */
gid_t      st_gid;          /* group ID of owner */
dev_t      st_rdev;         /* device ID (if special file) */
off_t      st_size;         /* total size, in bytes */
blksize_t  st_blksize;      /* blocksize for filesystem I/O */
blkcnt_t   st_blocks;       /* number of 512B blocks allocated */

time_t st_atime; /* time of last access */
time_t st_mtime; /* time of last modification */
time_t st_ctime; /* time of last status change */
};
```

# time and strace

- **time** [*options*] *command* [*arguments...*]
- Output:
  - real 0m4.866s: elapsed time as read from a wall clock
  - user 0m0.001s: the CPU time used by your process
  - sys 0m0.021s: the CPU time used by the system on behalf of your process
- **strace**: intercepts and prints out system calls to stderr or to an output file
  - \$ strace -o strace\_output ./tr2b 'AB' 'XY' < input.txt
  - \$ strace -o strace\_output2 ./tr2u 'AB' 'XY' < input.txt

# Homework 7

- Recall Homework 5!
- Rewrite `sfrob` using system calls (`sfrobu`)
- `sfrobu` should behave like `sfrob` except:
  - If `stdin` is a regular file, it should initially allocate enough memory to hold all data in the file all at once
  - It outputs a line with the number of comparisons performed
- Functions you'll need: `read`, `write`, and `fstat`  
(read the man pages, e.g. `man -S 2 read`)

# Homework 7

- Measure differences in performance between `sfrob` and `sfrobu` using the `time` command
- Estimate the number of comparisons as a function of the number of input lines provided to `sfrobu`
- Write a shell script “`sfrobs`” that uses `tr` and the `sort` utility to perform the same overall operation as `sfrob`
- Encrypted input -> `tr` (decrypt) -> `sort` (sort decrypted text) -> `tr` (encrypt) -> encrypted output

# Lab

[web.cs.ucla.edu/classes/winter16/cs35L/assign/assign7.html](http://web.cs.ucla.edu/classes/winter16/cs35L/assign/assign7.html)