

CS 35L

Week 8

TA: Tomer Weiss
Feb-19-2016

goo.gl/u74lLp

Slides

Announcements

- Student presentations today:
 - Michael: A Fish May Hold the Key to More Efficient Wireless Networks
 - Mu-Te: W3C Launches Effort to Replace Passwords
 - Hsien-Chih: Google's Balloon-Powered Internet Ready for Carrier Testing

web.cs.ucla.edu/classes/winter16/cs35L/assign/assign10.html

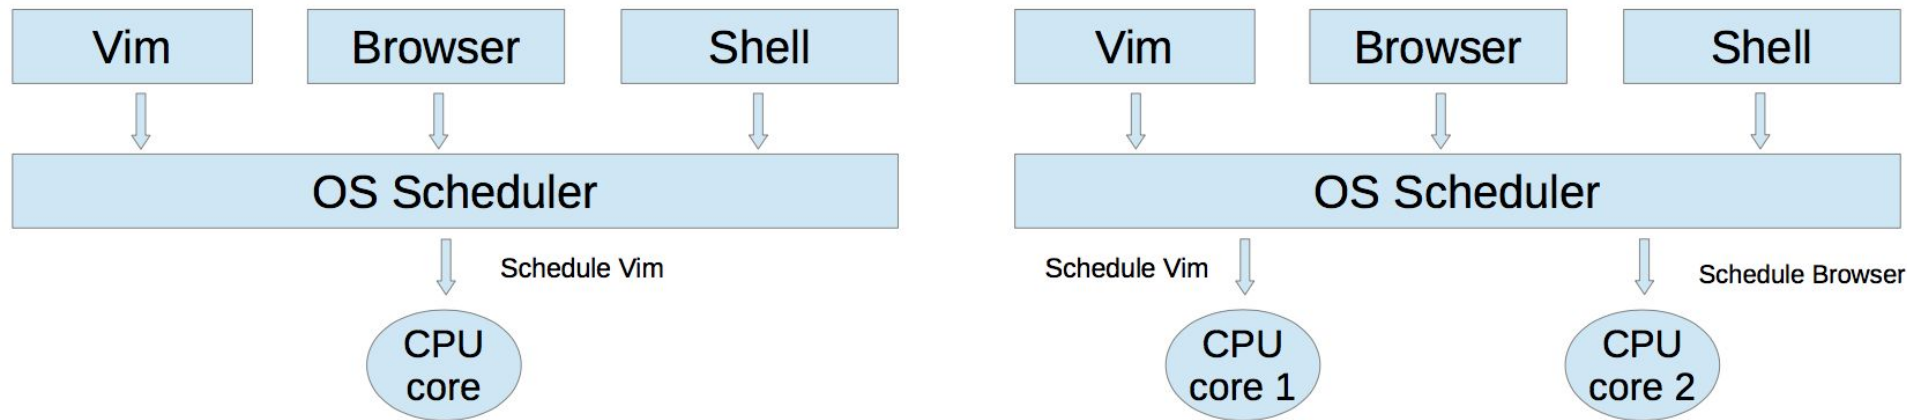
- Next week:
 - Write your topic [here](#)
 - Not registering your topic beforehand may result in rescheduling of your presentation
 - For reference on presentation, grading, please refer to this [rubric](#).

Multithreading/Parallel Processing

Week 8

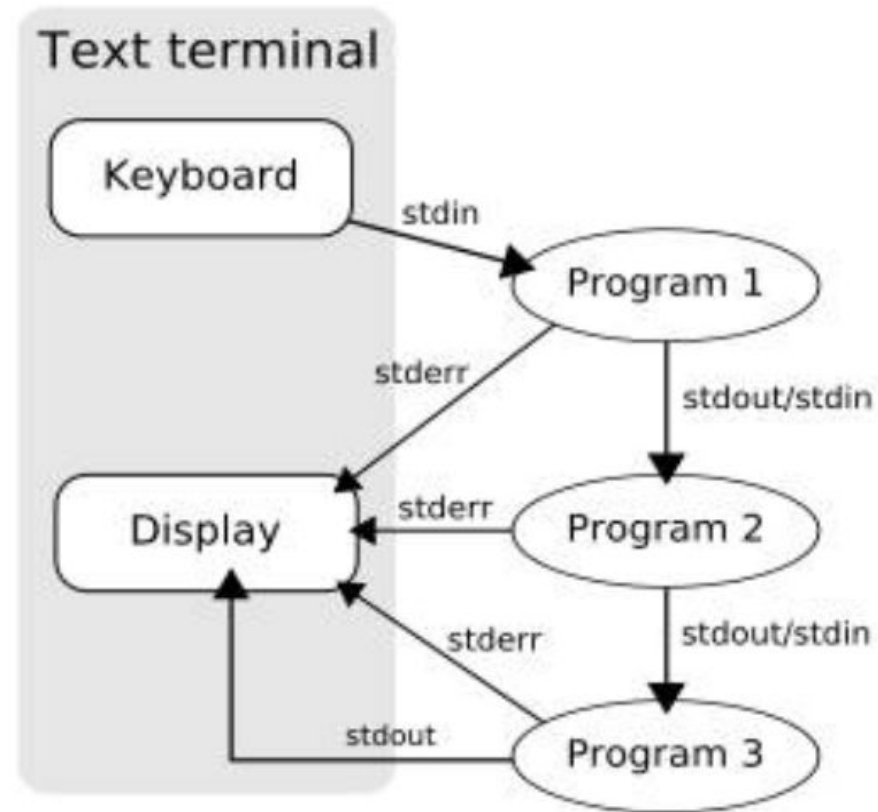
Multitasking

- Run multiple processes **simultaneously** to increase performance
- Processes do not share internal structures (stacks, globals, etc)
 - Communicate via **IPC** (inter-process communication) methods
 - Pipes, sockets, signals, message queues
- **Single core: Illusion** of parallelism by switching processes quickly (**time-sharing**). [Why is illusion good?](#)
- **Multi-core: True** parallelism. Multiple processes execute **concurrently** on different CPU cores



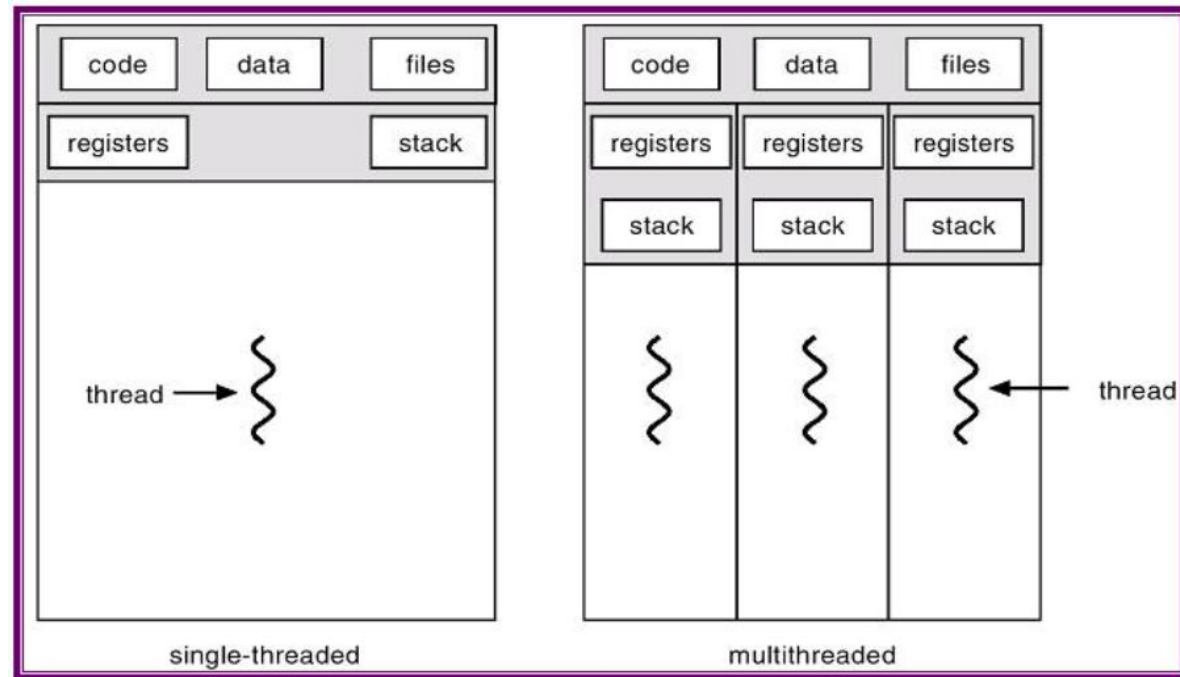
Multitasking

- `tr -s '[:space:]' '\n' | sort -u | comm -23 - words`
- Three separate processes spawned simultaneously
 - P1 - `tr`
 - P2 - `sort`
 - P3 - `comm`
- Common buffers (pipes) exist between 2 processes for communication
 - '`tr`' writes its `stdout` to a buffer that is read by '`sort`'
 - '`sort`' can execute, as and when data is available in the buffer
 - Similarly, a buffer is used for communicating between '`sort`' and '`comm`'



Threads

- A process can be
 - Single-threaded
 - Multi-threaded
- Threads in a process can run in parallel
- A thread is a lightweight process
- It is a basic unit of CPU utilization
- Each thread has its own:
 - Stack
 - Registers
 - Thread ID
- Each thread shares the following with other threads belonging to the same process
 - Code
 - Global Data
 - OS resources (files, I/O)



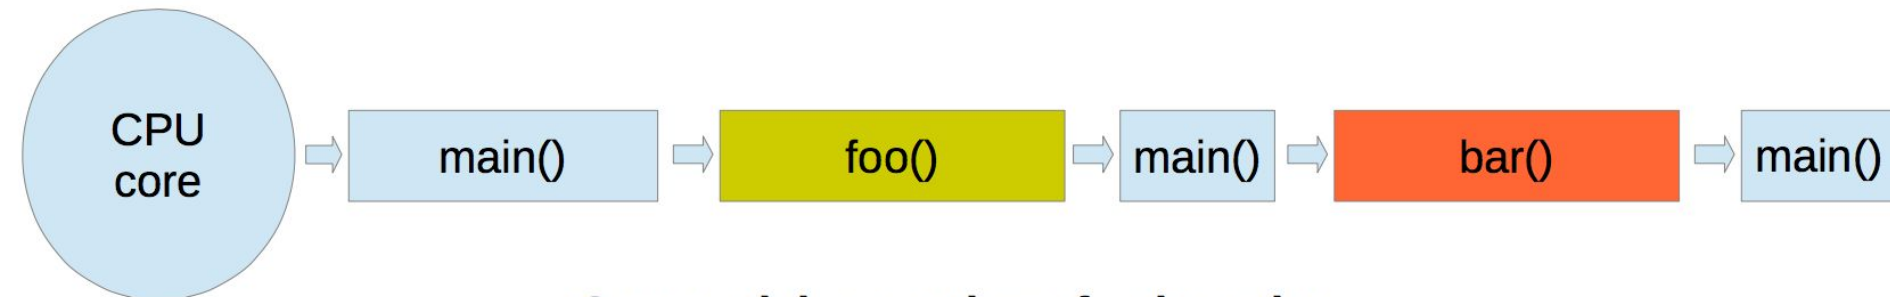
Single threaded execution

```
int global_counter = 0

int main()
{
    ...
    foo(arg1,arg2);
    bar(arg3,arg4,arg5);
    ...
    return 0;
}
```

```
void foo(arg1,arg2)
{
    //code for foo
}

void bar(arg3,arg4,arg5)
{
    //code for bar
}
```



Sequential execution of subroutines

Multi threaded execution (single core)

```
int global_counter = 0  
int main()  
{  
    ...  
    foo(arg1,arg2);  
    bar(arg3,arg4,arg5);  
    ...  
    return 0;  
}
```

```
void foo(arg1,arg2)  
{  
    //code for foo  
}  
void bar(arg3,arg4,arg5)  
{  
    //code for bar  
}
```




Time Sharing – Illusion of multithreaded parallelism
(Thread switching has less overhead compared to process switching)

Multi threaded execution (multiple cores)

```
int global_counter = 0  
int main()  
{  
    ...  
    foo(arg1,arg2);  
    bar(arg3,arg4,arg5);  
    ...  
    return 0;  
}
```

```
void foo(arg1,arg2)  
{  
    //code for foo  
}  
void bar(arg3,arg4,arg5)  
{  
    //code for bar  
}
```



True multithreaded parallelism

Multithreading properties

- Efficient way to **parallelize** tasks
- **Thread switches are less expensive** compared to process switches (context switching)
- Inter-thread communication is easy, via **shared global** data
- Need **synchronization** among threads accessing same data

Pthread API

```
#include <pthread.h>
```

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void* (*thread_function) (void*), void *arg);`
 - Returns 0 on success, otherwise returns non-zero number
- `void pthread_exit(void *retval);`
- `int pthread_join(pthread_t thread, void **retval);`
 - Returns 0 on success, otherwise returns non zero error number

```
#include<pthread.h> //Compile the following code as - gcc main.c -lpthread
#include<stdio.h>

void* ThreadFunction(void *arg) {
    long tID = (long)arg;
    printf("Inside thread function with ID = %ld\n", tID); pthread_exit(0);}

int main(int argc, char *argv[]) {
    const int nthreads = 5; pthread_t threadID[nthreads]; long t;
    for(t = 0; t < nthreads; ++t) {
        int rs = pthread_create(&threadID[t], 0, ThreadFunction, (void*)t);
        if(rs) {
            fprintf(stderr, "Error creating thread\n");
            return -1; }}
    printf("Main thread finished creating threads\n");
    for(t = 0; t < nthreads; ++t) {
        void *retVal;
        int rs = pthread_join(threadID[t], &retVal);
        if(rs) {
            fprintf(stderr, "Error joining thread\n");
            return -1;
        }}
    printf("Main thread finished execution!\n");
    return 0; }
```

Pthread API

Thread safety/synchronization

- **Thread safe function** - safe to be called by multiple threads at the same time. Function is free of 'race conditions' when called by multiple threads simultaneously.
- **Race condition** - the output depends on the order of execution
 - Shared data changed by 2 threads
 - `int balance = 1000`
 - Thread 1
 - T1 - read balance
 - T1 - Deduct 50 from balance
 - T1 - update balance with new value
 - Thread 2
 - T2 - read balance
 - T2 - add 150 to balance
 - T2 - update balance with new value

Thread safety/synchronization

- Order 1
 - balance = 1000
 - T1 - Read balance (1000)
 - T1 - Deduct 50
 - 950 in temporary result
 - T2 - read balance (1000)
 - T1 - update balance
 - balance is 950 at this point
 - T2 - add 150 to balance
 - 1150 in temporary result
 - T2 - update balance
 - balance is 1150 at this point
 - **The final value of balance is 1150**
- Order 2
 - balance = 1000
 - T1 - read balance (1000)
 - T2 - read balance (1000)
 - T2 - add 150 to balance
 - 1150 in temporary result
 - T1 - Deduct 50
 - 950 in temporary result
 - T2 - update balance
 - balance is 1150 at this point
 - T1 - update balance
 - balance is 950 at this point
 - **The final value of balance is 950**

Thread synchronization

- **Mutex (mutual exclusion)**
 - Thread 1
 - Mutex.lock()
 - Read balance
 - Deduct 50 from balance
 - Update balance with new value
 - Mutex.unlock()
 - Thread 2
 - Mutex.lock()
 - Read balance
 - Add 150 to balance
 - Update balance with new value
 - Mutex.unlock()
 - balance = 1100
- Only one thread will get the mutex. Other thread will **block in Mutex.lock()**
- Other thread can start execution only when the thread that holds the mutex calls **Mutex.unlock()**

Lab 8

- Evaluate the performance of multithreaded 'sort' command
 - `od -An -f -N 4000000 < /dev/urandom | tr -s ' ' '\n' > random.txt`
 - Might have to modify the command above
- Delete the empty line
 - `time -p sort -g --parallel=2 numbers.txt > /dev/null`
- Add /usr/local/cs/bin to PATH
 - `$ export PATH=/usr/local/cs/bin:$PATH`
- Generate a file containing 10M random **double-precision floating point numbers**, one per line with no white space
 - /dev/urandom: pseudo-random number generator

Lab 8

- `od`
 - write the contents of its input files to standard output in a user-specified format
 - Options
 - `-t f`: Double-precision floating point
 - `-N <count>`: Format no more than *count* bytes of input
- `sed, tr`
 - Remove address, delete spaces, add newlines between each float

Lab 8

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the `-g` option: compare by general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g file_name > /dev/null (1 thread)`
 - `$ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null`
 - Record the times and steps in `log.txt`

Lab

web.cs.ucla.edu/classes/winter16/cs35L/assign/assign8.html

Ray-Tracing

- **Powerful rendering technique in Computer Graphics**
- **Yields high quality rendering**
 - Suited for scences with complex light interactions
 - Visually realistic
 - Trace the path of light in the scene
- **Computationally expensive**
 - Not suited for real-time rendering (e.g. games)
 - Suited for rendering high quality pictures (e.g. movies)
- **Embarrassingly parallel**
 - Good candidate for **multi-threading**
 - Threads need **not synchronize** with each other, because each thread works on a different pixel

Ray-tracing






Image Source: POV Ray, Hall of Fame




Without ray tracing

With ray tracing



Ray-tracing

- Trace the path of a ray from the eye
 - **One ray per pixel** in the view window
 - The color of the ray is the color of the corresponding pixel
- Check for **intersection** of ray with scene objects.
- **Lighting**
 - **Flat shading** – The whole object has uniform brightness
 - **Lambertian shading** – Cosine of angle between surface normal and light direction



Homework 8

- Download the single-threaded ray tracer implementation
- Run it to get output image
- Multithread ray tracing
 - Modify main.c and Makefile
- Run the multithreaded version and compare resulting image with single-threaded one

Homework 8


- Build a multi-threaded version of Ray tracer
- Modify “main.c” & “Makefile”
 - Include <pthread.h> in “main.c”
 - Use “pthread_create” & “pthread_join” in “main.c”
 - Link with -lpthread flag (LDLIBS target)
- make clean check
 - Outputs “1-test.ppm”
 - Can see “1-test.ppm”
 - sudo apt-get install gimp (Ubuntu)
 - X forwarding (lnxsrvt)
 - gimp 1-test.ppm

1-test.ppm



Figure. 1-test.ppm

Homework 8 - antialiasing



Lab

web.cs.ucla.edu/classes/winter16/cs35L/assign/assign8.html